

Coreference Resolution without Global Attention via Quotient Graph Inference

Masum Billal

BILLALMASUM93@GMAIL.COM

Editor: TBD

Abstract

Modern approaches implicitly assume that global coreference resolution requires quadratic attention over the entire text sequence. This paper explores this unwritten assumption and examines whether global token-level cross-attention can be circumvented when features are factorized into distinct, architecturally separated representation channels. Our results indicate that quadratic token interaction can be restricted to localized context windows without severe performance degradation when combined with factorized feature channels, provided that connections are re-established between clusters among different windows. Long-range entity structure is subsequently recovered through a Graph Neural Network running over a compressed quotient graph of locally induced equivalence classes, operating at a true quadratic cost $\mathcal{O}(C^2)$ over $C \ll M$ clusters (where M is the number of mentions). Our core contribution is this structural decomposition principle that allows us to avoid training a single model to learn all features jointly. In our implementation, we use independent universal models as a realistic source of these distinct feature channels, while all inference is performed by lightweight learned components containing 16.1M parameters. Evaluated on the CoNLL-2012 benchmark with gold mentions, our approach achieves a baseline CoNLL F_1 score of 0.8646 on the test split, scaling to 0.8651 with validation-tuned decode-time merge thresholds. This indicates that strong document-level coreference performance can be achieved without global text-level attention through distinct channel factorization and structural quotient-graph inference.

Keywords: coreference resolution, representation factorization, quotient graph, Graph Neural Networks, window-based inference, transitive closure, efficient inference

1 Introduction

Coreference resolution is defined as the problem of clustering a set of entity mentions $\mathcal{M} = \{m_1, \dots, m_n\}$ into groups that are represented by one *canonical entity*. Formally, this induces a global equivalence relation $\sim$ over these mentions, where $m_i \sim m_j$ if and only if they refer to the same canonical entity. The true target output of the system is a partition $\mathcal{E} = \mathcal{M} / \sim$. From this perspective, corefer-

ence can be viewed as a structured prediction problem whose objective is to recover a latent quotient structure rather than to independently classify mention pairs or model token-level interactions.

In this work, we explicitly adopt this equivalence-relation formulation as the primary object of inference and study how it can be recovered from localized computations over restricted subsets of $\mathcal{M}$. We propose viewing the problem as a two-level decomposition: (i) construction of partial equivalence relations under localized observation constraints, and (ii) recovery of the global partition via composition and closure over these local relations in a compressed representation space. This formulation separates local relational induction from global consistency enforcement, allowing long-range dependencies to be recovered through structured aggregation over an intermediate quotient graph of locally induced equivalence classes rather than direct global interaction over the full mention set.

This paper challenges the assumption that long-range entity coreference requires monolithic token contextualization. Our results suggest that a substantial portion of the apparent reliance on global sequence-wide attention may arise from representation entanglement—where semantic, contextual, and discourse features are melted into a single joint vector space. Factorizing the required global signal into distinct representation channels permits the isolation of quadratic sequence operations to tight local windows, shifting the burden of long-range consistency from token-level attention to structural inference over a compressed quotient graph of locally induced equivalence classes.

The target decision boundary explicitly relies on the joint interpretation of these un-fused coordinate streams. Furthermore, we demonstrate that external symbolic channels—such as discrete lexical match dimensions—must not be collapsed early into scalar additions or post-hoc decoding heuristics. When integrated non-linearly via vector concatenation directly into the relational pairing heads, these side-channels unlock significant downstream coordinate synergy. By replacing independent pairwise matching heuristics with an explicit Graph Neural Network (GNN) operating over the quotient space, global consistency and competitive performance are recovered at an absolute computational complexity class completely decoupled from the document’s total mention count.

The remainder of this paper is organized as follows. Section 2 reviews related work and contextualizes our contribution. Section 3 outlines the overall workflow of the two-stage pipeline. Section 4 develops the mathematical foundations of representation factorization. Sections 5 and 6 detail the technical implementation of the Stage A local window inference procedure and Stage B quotient-space matching modules respectively. Section ?? describes experimental setup, and Section 7 presents empirical results and analysis.

2 Related Work

Scaling coreference resolution to long documents has traditionally relied on mitigating the quadratic cost of sequence-wide self-attention using two primary techniques:

Modern architectures frequently bypass full self-attention by substituting localized or sparse attention matrices, such as sliding windows and block-diagonal grids (Caciularu et al., 2023). While these methods reduce memory consumption during representation extraction, they remain structurally bound to an end-to-end fine-tuning paradigm where entangled relational features must be propagated across deep sequence lengths. In contrast, our approach restricts quadratic calculations strictly to isolated local execution windows, offloading long-range tracking to an uncontextualized cluster graph layer.

To bypass large all-pairs mention scoring matrices, several pipelines implement coarse-to-fine filtering heuristics (Dobrovolskii, 2021; Bohnet et al., 2023) or recursive mention clustering hierarchies (Luo et al., 2025). While effective at pruning computational overhead, these systems are vulnerable to cascading errors during iterative merging steps. We depart from this by reformulating cross-window resolution as a formal completion task over a highly compressed quotient graph. By leveraging an un-fused coordinate basis rather than early scalar similarities, our method preserves fine-grained channel identities over extreme token distances without requiring global sequence contextualization.

3 Solution Overview

To bypass global quadratic text attention, the framework processes an incoming document through a two-stage hierarchical inference pipeline that operates first on bounded local context and then on a compressed quotient representation of induced entities.

1. **Decoupled Encoding Stream:** Each mention is mapped into two fixed, pretrained representation spaces capturing distinct semantic and contextual information. These encoders are frozen throughout training, and their outputs are never fused into a single joint representation at the token level. All subsequent reasoning is performed over these fixed, separated feature views.
2. **Stage A (Intra-Window Relational Construction):** The document is partitioned into a set of sentence-aligned windows of bounded size. Within each window, all pairwise mention interactions are computed using a local scoring function over the distinct representations. A transitive inference procedure (union-find) is then applied strictly within each window, producing a set of locally consistent equivalence structures.
3. **Quotient Representation Induction:** The output of Stage A is a collection of locally inferred equivalence groups, each representing a collapsed subset of

mentions. These groups are treated as atomic nodes for all subsequent computation, inducing a reduced representation space in which reasoning operates over clusters rather than raw mentions.

4. **Stage B (Cross-Window Quotient Matching via GNN):** Global entity consistency is recovered by evaluating structural interactions over the induced clusters in the form of an antecedent ranking task. A Graph Neural Network performs message passing over the quotient graph topology, processing size-invariant node features derived from local non-parametric cluster aggregations combined with concatenated lexical signals. Final document-level entities are obtained by computing the transitive closure over this pointer graph, yielding a globally consistent partition of mentions.

4 Representation Factorization Framework

Monolithic language models rely on end-to-end contextualization over full-document sequences to implicitly disentangle semantic, discourse, and referential dependencies. This couples representation learning with relational inference and induces unnecessary computational cost. We instead hypothesize that coreference can be reformulated as a structured inference problem over localized observations combined through a compressed global reasoning space. Under this view, global entity structure does not require direct token-level interaction, but can be recovered from compositions of locally induced relational evidence. This motivates a factorization of inference into distinct representation channels that preserve complementary signals for downstream relational reconstruction.

Assumption 1 (Distinguishable Feature Representation Basis) *Let $\mathcal{V}$ denote the complete decision space required to resolve entity identity. We assume that discriminative identity information can be systematically recovered across d distinct, architecturally separated feature channels whose signals remain structurally intact when preserved separately:*

$$\mathcal{V} = \mathcal{V}_{context} \times \mathcal{V}_{semantic} \times \dots$$

where individual channels preserve specialized signals that are obscured when compressed into low-dimensional similarity measures or early aggregation operations.

To build this basis, we incorporate a dual-stream architecture tracking context and semantics independently. For a given mention m_i , a contextual vector is derived from a frozen RoBERTa-large backbone (the concatenated start/end subtoken states, $\mathbb{R}^{2048}$), and a semantic representation is computed using a frozen BGE-large model ($\mathbb{R}^{1024}$). Each frozen channel is linearly projected into an isolated low-rank coordinate path and concatenated into the factorized mention representation:

$$\mathbf{h}_i = [P_{ctx}(f_{RoBERTa}(m_i)) \parallel P_{bge}(f_{BGE}(m_i))]$$

The local Stage A scorer projects each channel to 1024 dimensions ($\mathbf{h}_i \in \mathbb{R}^{2048}$), while the cluster-level Stage B uses a tighter 512-dimensional projection (Section 6), reflecting its operation over pooled clusters rather than raw mentions. This projection enforces a low-rank bottleneck that regularizes the interaction head; we find the width near-optimal: removing the projection entirely (training the head on raw concatenated features) overfits and degrades end-to-end performance by -1.58 , while widening it yields no gain, consistent with the coreference signal being low-rank as established below.

To isolate whether this specific factorization drives system capability, we introduce a channel control configuration $\text{channel} \in \{\text{both}, \text{ctx}, \text{bge}\}$. When a single channel is isolated, the unused projection layer is dropped entirely and the upstream scoring heads are dynamically resized. This architecture guarantees that single-channel variants are trained completely from scratch rather than evaluating zeroed vectors at inference time, avoiding confounding the capacity of the model head with the signal density of the channel.

The structural necessity of this factorization is verified through single-channel control experiments. Isolation of the contextual channel ($\text{channel} = \text{ctx}$) drops full global performance to $0.8416 F_1$, while isolating the semantic channel ($\text{channel} = \text{bge}$) triggers a complete collapse to $0.7813 F_1$. This confirms that the target decision boundary explicitly relies on the joint interpretation of un-fused coordinate streams. Pruning either single coordinate path prevents the model from mapping cross-window dependencies effectively.

The critical property of this system is not that the underlying encoders are independent, but that the decision information is supplied through separate coordinate systems whose signals remain explicitly distinguishable throughout inference. Frozen universal encoders provide one convenient realization of this condition; any mechanism capable of producing comparably distinct channels would satisfy the same architectural requirement. The contribution of this work is therefore not the use of frozen models itself, but the demonstration that preserving factorized information channels can mitigate the reliance on global token-level attention within document-level coreference resolution.

This factorization is explicitly validated by an empirical probe: computing raw cosine similarities over these embedding dimensions directly yields an AUC of ≈ 0.57 on separating coreferent from non-coreferent pairs. Conversely, a shallow learned linear probe trained over the un-fused coordinate vectors $(\mathbf{h}_i, \mathbf{h}_j)$ separates identity at an AUC of 0.91 . This demonstrates that substantial coreference information is present in the preserved multi-channel representation and remains linearly recoverable. The large gap between cosine similarity and linear-probe performance suggests that scalar similarity measures discard discriminative structure that remains available when the separate coordinate channels are preserved.

5 Stage A: Bounded Local Equivalence Reconstruction

Given the factorized representations $\mathbf{h}_i$, we restrict initial quadratic interaction to bounded observation domains to evaluate the structural locality of coreference decisions.

Assumption 2 (Transitive Link Reconstruction) *Global entity paths are locally dense: any two coreferent mentions $m_a, m_b \in \mathcal{M}$ are connected via a finite chain of intermediate local links, where each adjacent link pair falls within a bounded sequence window.*

The architectural viability of restricting token-level quadratic attention to isolated local windows is supported by the factorization established in Section 4. Because the coordinate representations retain un-collapsed channel signals, local relational boundaries preserve substantial identity information that can drop under conventional scalarized similarity spaces, mitigating the contextual degradation typically observed in naive windowed models.

Rather than imposing a rigid token-block partition, we implement a sentence-aligned greedy windowing strategy. The input document sequence is grouped into disjoint context windows W_k of at most $K = 256$ subtokens by greedily packing intact, whole sentences. A sentence is appended to the current active execution block only if its subtoken sequence fits within the $K = 256$ bound; otherwise, a new window is opened. This design ensures that no sentence boundary is ever fractured across a window interface, presenting the transformer backbone with complete context. Each mention is assigned to exactly one window via a mapping function $\pi : \mathcal{M} \rightarrow \mathcal{W}$.

For every pair of mentions co-located within W_k , an antecedent scoring head (**AntecedentScorer**) calculates local compatibility alongside bucketed distance tokens:

$$s(i, j) = \text{MLP}([\mathbf{h}_i \parallel \mathbf{h}_j \parallel \mathbf{e}_{\text{dist}(i-j)}])$$

A local equivalence relation $\sim_k$ is induced via thresholding against a learned null bias parameter τ . Computing the transitive closure via a union-find algorithm strictly within each window boundaries yields a set of local equivalence classes:

$$\mathcal{C}_k = W_k / \sim_k^*$$

The union of these sub-clusters across all windows forms our complete intermediate entity space $\mathcal{C} = \bigcup_{k=1}^W \mathcal{C}_k$. Elements of $\mathcal{C}_k$ constitute local connected components under localized constraints, rather than a final global quotient space. Each local cluster $C \in \mathcal{C}$ is subsequently treated as an atomic node in a quotient graph of locally induced equivalence classes for cross-window inference layers.

Evaluating the isolated performance of Stage A against window-split gold targets (crediting locally achievable coreference) yields a baseline score of 0.8948 F₁ on

CoNLL-2012. When these same window-local clusters are left unmerged and evaluated directly against full-document gold targets, Stage A achieves a score of 0.7439 F_1 . Crucially, local equivalence reconstruction alone retains the vast majority of final performance without requiring global sequence-wide matching dependencies, trailing the complete system by 0.1207 points. This provides a central empirical finding of this work: a large fraction of document-level coreference structure remains highly concentrated within bounded local neighborhoods.

6 Stage B: Quotient-Space Completion via Graph Neural Networks

The intermediate local equivalence classes derived in Section 5 act as atomic nodes in a highly compressed graph. Rather than executing unconstrained mention-pair matching across window splits, Stage B is formulated as an antecedent ranking problem over the quotient graph topology. The objective is to recover global coreference consistency by establishing a directed cross-window pointer forest over macro-nodes, bounded by a strict complexity class of $\mathcal{O}(M + C^2)$ where $C \ll M$.

Proposition 1 *Let $\mathcal{C} = \bigcup_k \mathcal{C}_k$ denote the set of locally induced equivalence classes where $\mathcal{C}_k = W_k / \sim_k^*$. If an operator $\approx$ correctly identifies cross-window equivalences between elements of $\mathcal{C}$, then the global entity partition $\mathcal{E} = \mathcal{M} / \sim$ is recovered exactly by the transitive closure of $\approx$ lifted back onto $\mathcal{M}$.*

6.1 Quotient Graph Formulation & Node Pooling

Let $\mathcal{C} = \{C_1, C_2, \dots, C_C\}$ be the set of clusters extracted by Stage A across all sequence windows. Each member mention is re-projected into Stage B’s 512-dimensional channel paths and concatenated to $\tilde{\mathbf{h}}_i \in \mathbb{R}^{1024}$. To map a cluster into the macro-graph without smearing discriminative features, we apply a size-normalized log-mean-exp pooling operator (`lse`) over its members, followed by a learned input layer producing the base node representation $\mathbf{v}_c \in \mathbb{R}^{512}$:

$$\mathbf{v}_c = W_{\text{in}} \log \left(\frac{1}{|C_c|} \sum_{m_i \in C_c} \exp(\tilde{\mathbf{h}}_i) \right)$$

This soft-maximum operator preserves the vector identity of the single most discriminative mention within the sub-cluster (e.g., a distinct proper noun surface string), bypassing the signature attenuation or “smearing” penalty typical of uniform attention pooling mechanisms.

6.2 Structural Message Passing & Outer Residuals

The pooled cluster nodes are contextualized globally using a 2-layer Transformer encoder serving as a fully connected quotient message-passing network. To prevent *all-null collapse*—where deep transformer layers over-smooth representations

to near-identical vectors (mean pairwise cosine $\rightarrow 1.000$) and completely erase the discriminative features required to resolve distinct entity identities—we enforce a rigid outer residual connection:

$$\mathbf{H} = \mathbf{H}_0 + \text{GNN}(\mathbf{H}_0)$$

where $\mathbf{H}_0$ represents the initial array of un-smoothed cluster node vectors augmented with window position embeddings. The residual path ensures that localized discriminative structures always reach the final scoring heads, shifting the role of graph message passing from base representation creation to the relational refinement of existing entity boundaries.

6.3 Concatenated Lexical Integration & Decode Tuning

To assist the neural coordinates across vast token spans, symbolic side-channels are integrated directly into the pair interaction MLP. Crucially, experimental validation reveals that collapsing discrete lexical matching features into an additive scalar penalty or a post-hoc decoding heuristic creates threshold coupling that degrades performance on non-lexical categories. We instead concatenate the symbolic signals directly into the joint neural feature space, where the pairwise head consumes the contextualized node streams alongside the lexical vectors:

$$\mathbf{x}_{\text{pair}} = [\mathbf{v}_A \parallel \mathbf{v}_B \parallel \mathbf{v}_A \odot \mathbf{v}_B \parallel |\mathbf{v}_A - \mathbf{v}_B| \parallel \mathbf{x}_{\text{lex}}] \in \mathbb{R}^{2051}$$

where $\mathbf{x}_{\text{lex}} \in \mathbb{R}^3$ tracks:

1. **IDF-weighted token Jaccard similarity:** $\sum_{t \in A \cap B} \text{idf}(t) / \sum_{t \in A \cup B} \text{idf}(t)$, where document-local $\text{idf}(t) = \log(M / (1 + \text{df}(t))) + 1$.
2. **Substring containment:** A binary indicator marking if any content mention (excluding bare pronouns) in A is a substring of any content mention in B , or vice versa.
3. **Exact mention match:** A binary flag tracking if the two clusters share an identical content mention surface text string.

This concatenation allows the relational head to interpret symbolic signals non-linearly and conditionally based on the surrounding neural context, unlocking a system baseline performance of 0.8646 F_1 on the CoNLL-2012 test split.

Final entity chains are recovered by allowing each macro-node to predict a single antecedent pointer among strictly earlier temporal windows, thresholded against a learned null node τ . By implementing a post-hoc, validation-tuned calibration sweep over the decode-time merge threshold ($\delta = +0.50$), the final system successfully mitigates residual cross-window fragmentation and reaches our peak benchmark performance of **0.8651** F_1 .

7 Results and Analysis

All evaluation experiments are conducted on the test split of the CoNLL-2012 English coreference benchmark (Pradhan et al., 2012) utilizing gold mention boundaries. Following empirical data mixture optimizations, our lightweight scoring modules are trained on a uniform corpus mixture consisting of CoNLL-2012, LitBank, CorefUD, and a subsampled partition of 8,000 documents from PreCo. Training enforces a uniform loss across all datasets without negative window-pair subsampling to preserve distributional balance. All baseline results cited from prior literature correspond strictly to their respective gold-mention evaluation configurations.

We isolated the performance of individual channels against the full dual-channel baseline. All variations utilize an identical sentence-aligned window configuration ($K \leq 256$) and the Cluster-GNN architecture trained on the uniform multi-corpus mixture. Table 1 details the results.

Channel Configuration	Stage A	Stage A+B	Δ
Both Channels (BGE + RoBERTa)	0.8948	0.8646	—
RoBERTa Contextual Only (ctx)	0.8727	0.8416	− 2.30
BGE Semantic Only (bge)	0.8093	0.7813	− 8.33

Table 1: Channel ablation control results (CoNLL F_1) on CoNLL-2012 test data splits demonstrating non-redundant coordinate paths.

The empirical trajectory reveals a clear asymmetry between the two coordinate views. The contextual channel (ctx) is heavily dominant, achieving 0.8416 F_1 on its own—within 2.30 points of the complete system. Conversely, the semantic channel (bge) drops significantly to 0.7813, showing that cross-window cluster bridging retains very little tracking signal when divorced from contextual sequence history.

Granular cluster analysis uncovers a strict architectural division of labor between the channels. Pronominal groups (PRON-only) are almost completely dependent on the contextual channel, where the ctx-only run reaches an integrated baseline score of 0.6556 end-to-end, compared to the semantic bge-only run which collapses to 0.5173. Conversely, proper nouns (PROPN-only) lean securely on semantic assignments, where bge-only maintains a stable footing (0.7276) relative to context alone (0.7983). This split confirms that pronouns depend almost entirely on contextual signals, while surface variations of proper nouns rely on global text-invariant semantic spaces.

This division of labor is governed by a conditional-independence criterion that we sharpen with a direct linear-dependence probe. A morphological agreement profile—the gender, number, and person of a pronoun—is recovered from the frozen RoBERTa contextual space by a *linear* probe at 96%–100% accuracy, confirming that it is a linearly-dependent function of a channel already present rather than an

independent signal. A feature that is linearly recoverable from an existing channel carries no complementary information for a head that already reads that channel. Preserving factorized channels is therefore beneficial only when those channels are conditionally independent given the label, not merely jointly correlated with it—a criterion that predicts both which channels earn their place and which are redundant.

7.1 Benchmarks

Table 2 establishes the primary system performance evaluated against state-of-the-art architectures using identical gold mention contexts.

Approach	CoNLL	MUC	B ³	CEAF _e
Stage A Only (No Merge Baseline, Full-Doc Gold)	0.7439	0.8460	0.6910	0.6950
Stage A Only (Window-Split Gold Baseline)	0.8948	0.9272	0.8889	0.8684
Stage A + B GNN ($\mathcal{O}(C^2)$)	0.8646	0.9247	0.8437	0.8254
Stage A + B GNN + Decode Tuning (δ)	0.8651	0.9250	0.8441	0.8259
Dobrovolskii (2021)	0.8180	0.8570	0.8000	0.7970
Caciularu et al. (2023)	0.8360	0.8740	0.8200	0.8130
Bohnet et al. (2023)	0.8380	0.8760	0.8220	0.8150
Luo et al. (2025)	0.8510	0.8890	0.8350	0.8290

Table 2: Full-document CoNLL-2012 test results on **gold mention boundaries**. Baseline systems are mapped strictly to their gold-mention settings to isolate linking capability.

Cluster Category Type	CoNLL	MUC	B ³	CEAF _e
PROPN-only	0.8697	0.8708	0.8671	0.8710
NOUN-only	0.7793	0.7827	0.7798	0.7753
PRON+NOUN	0.7573	0.7389	0.7354	0.7975
PRON+PROPN	0.7468	0.7644	0.7190	0.7571
all-mixed	0.7137	0.7678	0.6781	0.6953
NOUN+PROPN	0.7078	0.7102	0.6977	0.7155
PRON-only	0.6861	0.7215	0.6636	0.6731

Table 3: Definitive per-cluster category breakdown for the baseline Stage A + B GNN pipeline, CoNLL-2012 test.

7.2 Complexity and Operational Scale

Since window size is fixed, total sequence encoding complexity runs at $\mathcal{O}(WK^2) = \mathcal{O}(NK)$, avoiding full-sequence token quadratics. Within Stage A, localized mention scoring operates at $\mathcal{O}(W \cdot (M/W)^2) = \mathcal{O}(M^2K/N)$.

For Stage B cross-window matching, the integration of macro-node pooling maps the trainable computation path directly into the number of locally induced sub-clusters:

$$T_{\text{Stage B}}(C) = \mathcal{O}(M) + \mathcal{O}(C^2)$$

where C represents the total number of clusters generated by Stage A per document. Because $C \ll M \ll N$, the absolute computation budget required for global entity synthesis is entirely decoupled from the quadratic scale of the document sequence or token length.

Space Layer	Graph Nodes	Evaluation Points	Algorithmic Mode
Token Sequence Space	617.4	381,183	$\mathcal{O}(NK)$ Windowed Attention
Mention Graph Space	57.1	3,260	Unwindowed Mention Baseline
Quotient Graph Space	20.7	428	$\mathcal{O}(C^2)$ Trainable Graph Decisions

Table 4: Average document graph dimensions and operational scales evaluated on CoNLL-2012 test data splits.

Empirical metrics collected on the CoNLL-2012 test set (Table 4) validate this structural optimization. Mean scales establish document dimensions of $N = 617.4$ tokens, $M = 57.1$ nominal mentions, and $C = 20.7$ locally induced macro-nodes. Stage A contracts the entity node space down to 36.3% of its uncollapsed volume, yielding a 2.76:1 node compression ratio.

This contraction drastically reduces the learned evaluation space. While an unwindowed all-pairs system must parameterize and score 3,260 potential mention edges, Stage B operates over a highly compressed topology of only 428 candidate macro-node interactions. This represents an empirical **7.6 $\times$ structural reduction** in the learned decision graph size. By handling the internal `lse` member pooling as a non-parametric aggregation over cached coordinate segments, the parameter execution path is restricted to a compact macro-node topology.

Our experimental findings indicate that the explicit core theoretical assumptions are empirically viable:

1. **Preservation of Feature Distinctness (Assumption 1):** Reaching a baseline F_1 score of 0.8646 is consistent with the premise that preserving distinguishable feature channels retains structural identity information that remains accessible for downstream reconstruction. This observation is supported by the notable performance gap between raw cosine evaluation and the un-fused multi-channel linear probe.
2. **Validation of Locality-Transitivity (Assumption 2 & Proposition 1):** Stage B achieves stable macro-transitive lifting across window boundaries, supported by an average fragmentation profile of only 1.45 nodes per true entity.

By locking the primary encoder backbones and routing cross-window connections through the compressed quotient graph, the model significantly reduces computational overhead. As itemized in Table 5, trained parameters account for only 2.28% (16.1M) of the total 706.13M parameter budget.

Component	Parameter Count
Stage A (AntecedentScorer)	8.43M
Stage B (ClusterGNN)	7.70M
RoBERTa-large	355.00M
BGE-large	335.00M
Total Parameters	706.13M

Table 5: System parameter budget distribution.

Architecture Model	Inference Speed	Incremental Allocated	Total Reserved VRAM
Quotient-GNN ($\mathcal{O}(C^2)$)	19.4 ms / 1k tok	64.2 MB	2.41 GB
Unwindowed Baseline ($\mathcal{O}(M^2)$)	13.8 ms / 1k tok	198.5 MB	2.54 GB

Table 6: Runtime and memory profile measurements collected per-document during inference on a single RTX 4060 Laptop GPU.

We note an explicit architectural distinction regarding our computational profile: the primary claim of this work is not that the overall pipeline contains zero global-attention calculations, as the frozen backbones utilize internal transformer layers. Rather, the claim demonstrates that **task-specific entity inference and downstream cluster assembly are executed completely without document-wide token interaction**.

In inference, this architectural factorization limits memory usage. The Quotient-GNN restricts peak active incremental allocations to 64.2 MB by preventing the

materialization of the full unwindowed mention-pair matrix, yielding a $3.1\times$ optimization in active running footprint over an unwindowed setup. Total cumulative VRAM footprint remains bounded near 2.41 GB due to the shared memory overhead of the frozen token embedding caches (200 MB) and CUDA driver baselines.

An honest evaluation of wall-clock times reveals that at standard lengths (mean 617.4 subtokens), the unwindowed architecture retains a minor speed advantage (13.8 ms vs 19.4 ms per document). Because sequence spans are short in this corpus, unwindowed $\mathcal{O}(N^2)$ mechanics are not yet computationally punitive, allowing an unwindowed layout to benefit from highly parallel vectorized GPU paths. The windowed system encounters a minor Python execution loop penalty when processing small windowed batches sequentially. The asymptotic runtime advantage of the $\mathcal{O}(NK)$ design manifests on significantly longer document horizons, while the memory runtime ceiling reduction remains immediately visible.

8 Conclusion

Our findings demonstrate that global entity resolution can be successfully recovered from local inference paired with quotient-graph reconstruction. Our results suggest that document-level coreference can be performed without global token-level attention when entity inference is reformulated as local equivalence reconstruction followed by quotient-space completion. In our implementation, preserving separate representation channels was sufficient to recover global entity structure through local edge induction and quotient-space reconstruction without requiring document-wide token attention.

These results raise the possibility that some relational prediction problems may admit similar structural decompositions into localized inference and quotient-space reconstruction, particularly where decision boundaries can be expressed as stable compositions of distinguishable information channels. In the setting studied here, exhaustive sequence-wide cross-attention does not appear to be an absolute prerequisite for strong document-level coreference performance due to the inherent equivalence relation structure. Under this decentralized processing paradigm, downstream multi-task architectures could allocate parameter capacity toward isolating distinct specialized feature views, leveraging a lightweight central language infrastructure purely to execute structural orchestrations and macro-equivalence pooling rather than computing dense all-to-all token interactions over an entire contextual window.

References

Bernd Bohnet, Chris Alberti, and Michael Collins. Coreference resolution through a seq2seq transition-based system. *Transactions of the Association for Computational Linguistics*, 11:212–226, 2023.

- Avi Caciularu, Jacob Goldberger, and Ido Dagan. Peek across sentences for knowledge-grounded coreference resolution. In *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics (ACL)*, pages 4078–4094, 2023.
- Vladimir Dobrovolskii. Word-level coreference resolution. In *Proceedings of the 2021 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, pages 7670–7675, 2021.
- Ke Luo, Yu Bai, Shuai Si, Chao Gao, Zhen Wang, Yelong Shen, Weizhu Li, Zhiyuan Liu, Yanyan Han, Jian Wu, Chuanyi Kong, and Maosong Sun. ImCoref-CeS: An improved lightweight pipeline for coreference resolution with LLM-based checker-splitter refinement, 2025. Preprint.
- Sameer Pradhan, Alessandro Moschitti, Nianwen Xue, Olga Uryupina, and Yuchen Zhang. CoNLL-2012 shared task: Modeling multilingual unrestricted coreference in OntoNotes. In *Joint Conference on EMNLP and CoNLL — Shared Task*, pages 1–40, 2012.